

Disposition 11: State Machine Replication

Mathias Ravn Tversted

January 6, 2020

Table of contents

State machine

Replicated state machine

Consistency: TOB

- Synchronous implementation

- Async implementation

 - Enten den her

 - Eller den her

Client-centric consistency

Definition of a state machine

A state machine consists of

- ▶ A set of states
- ▶ Initial state $State_0 \in States$
- ▶ Set of *Inputs*
- ▶ Set of *Outputs*
- ▶ Transition function
 $T : States \times Inputs \rightarrow States \times Outputs$

IF and definition of Replicated State Machine

We have I/O ports RSM_i and the following ports

- ▶ $\forall i Received_i$ outputs what IF has Received
- ▶ $Process$ says what should be processed next
- ▶ $Deliver_i$ instructs IF to deliver the next message to S_i

It has the following safety requirements:

- ▶ **Validity**: If honest server outputs (y_1, \dots, y_n)
Then $\exists (x_1, \dots, x_n) : (y_1, \dots, y_n) = M(x_1, \dots, x_n)$
- ▶ **Agreement**: If honest server outputs (y_1, \dots, y_n)
then all other servers output at least some prefix of that, or vice versa

Ideal Functionality of an RSM

- ▶ Let $State = State_0$. Foreach RSM_i , $Q_i = \emptyset$.
- ▶ Q_i is the outputs for S_i , which has not yet been delivered. Let $UnProcessed = \emptyset$
- ▶ On input x to RSM_i , output x on $Received_i$, add x to $UnProcessed$
- ▶ On input x on $Process$. If $x \in Unprocessed$, compute $(State', y) = T(state, x)$. Add y to Q_i . Pop x from $UnProcessed$
- ▶ On input $Deliver_i$, where $Q_i \neq \emptyset$, remove first $y_i \in Q_i$ and output y on RSM_i

Consistency

In order to keep everyone consistent, we need to build it on *Total-Ordered Broadcast*. Because we have state machines, the order of processing matters, therefore we need to be able to have all the machines process things in the same order. These can be

- ▶ Synchronous TOB
- ▶ Async TOB
- ▶ Eventually synchronous broadcast

Synchronous Implementation of TOB

- ▶ On x at S_i , flood it
- ▶ All servers S_i keep *UnQueued* _{i} of received messages. These could be received in arbitrary order
- ▶ All servers S_i keep a set *Queued* _{i} . When they move messages from *UnQueued* to *Queued*, they do so in the same order
- ▶ There is a leader L (sequencer), who makes the order. Corrupted leader may break liveness but not safety
- ▶ If a message is put into *Queued*, they do it in the same order. Leadership goes round-robin. This guarantees liveness
- ▶ This happens with blocks

Adapting IF to Async broadcast

Now we have a notion of an epoch, *Unqueued* and we transmit multiple inputs in combined blocks.

1. On input x to TOB, add it to *Unqueued*
2. Leader of epoch is $P_i \bmod n$
3. Leader adds *Unqueued* to block and broadcasts it
4. When receiving block, add it to Q

Core-set selection: EKSTRA STUFF DER KAN SPRINGES OVER?

We cannot reliably wait for the leader in each epoch. So we have everyone propose a block, and when a block is *seen by many honest*, we can trust it. We simply have everyone collect these blocks and take the union of them. We have Byzantine Agreement available to us in the async model.

Tidsmæssigt problematisk: Core-set selection

Per fig. 11.4

Client vs. Server side consistency

Server-side, which is all honest servers executing commands in the same order

Client-side, which is servers being behind other servers. This is why agree requires that for honest servers, that their output is a prefix of another honest parties output